

Assignment 4

Question NO. 1 Using sklearn.datasets.load_diabetes apply Variance method for removing the constant column also after applying the Variance method apply multi linear regression on that data

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_diabetes
from sklearn.feature_selection import VarianceThreshold
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.metrics import accuracy_score

diab = load_diabetes()

X = diab.data
y = diab.target

diab.feature_names

['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']

diab_data = pd.DataFrame(diab.data, columns=diab.feature_names)
diab_data.head()
```

	age	sex	bmi	bp	s1	s2
s3 \						
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821
	0.043401					
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163
	0.074412					
2	0.085299	0.050680	0.044451	-0.005671	-0.045599	-0.034194
	0.032356					
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991
	0.036038					
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596
	0.008142					

	s4	s5	s6
0	-0.002592	0.019908	-0.017646
1	-0.039493	-0.068330	-0.092204
2	-0.002592	0.002864	-0.025930
3	0.034309	0.022692	-0.009362
4	-0.002592	-0.031991	-0.046641

```
varT = VarianceThreshold()

X_filt = varT.fit_transform(X)
```

```

X_filt
array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.00259226,
         0.01990842, -0.01764613],
       [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
        -0.06832974, -0.09220405],
       [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
        0.00286377, -0.02593034],
       ...,
       [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
        -0.04687948,  0.01549073],
       [-0.04547248, -0.04464164,  0.03906215, ...,  0.02655962,
        0.04452837, -0.02593034],
       [-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,
        -0.00421986,  0.00306441]])

op = varT.fit(X)
op.get_support()

array([ True,  True,  True,  True,  True,  True,  True,  True,  True,
        True])

X_train,X_test,y_train,y_test = train_test_split(X_filt,y, test_size
=0.25,random_state = 17)
print(X_train.shape,X_test.shape)
print(y_train.shape,y_test.shape)
model = LinearRegression()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print("R2 Score ",r2_score(y_test,y_pred))

(331, 10) (111, 10)
(331,) (111,)
R2 Score    0.49062912872854303

```

Question NO.2 Using sklearn.datasets.load_wine Apply Correlation and make a heat map using seaborn and remove the highly correlated columns if exist and the apply SVM and get the best accuracy by changing the Hyperparameters

```

from sklearn.datasets import load_wine

wine = load_wine()

X2, y2 = wine.data,wine.target
wine['target'] =y2

wine.feature_names

['alcohol',
 'malic_acid',
 'ash',
 'alcalinity_of_ash',

```

```

'magnesium',
'total_phenols',
'flavanoids',
'nonflavanoid_phenols',
'proanthocyanins',
'color_intensity',
'hue',
'od280/od315_of_diluted_wines',
'proline']

```

```
wine_df = pd.DataFrame(wine.data, columns=wine.feature_names)
```

```

corr_matrix = wine_df.corr()
corr_matrix

```

	alcohol	malic_acid	ash	\
alcohol	1.000000	0.094397	0.211545	
malic_acid	0.094397	1.000000	0.164045	
ash	0.211545	0.164045	1.000000	
alcalinity_of_ash	-0.310235	0.288500	0.443367	
magnesium	0.270798	-0.054575	0.286587	
total_phenols	0.289101	-0.335167	0.128980	
flavanoids	0.236815	-0.411007	0.115077	
nonflavanoid_phenols	-0.155929	0.292977	0.186230	
proanthocyanins	0.136698	-0.220746	0.009652	
color_intensity	0.546364	0.248985	0.258887	
hue	-0.071747	-0.561296	-0.074667	
od280/od315_of_diluted_wines	0.072343	-0.368710	0.003911	
proline	0.643720	-0.192011	0.223626	

	alcalinity_of_ash	magnesium	
total_phenols \			
alcohol	-0.310235	0.270798	
0.289101			
malic_acid	0.288500	-0.054575	-
0.335167			
ash	0.443367	0.286587	
0.128980			
alcalinity_of_ash	1.000000	-0.083333	-
0.321113			
magnesium	-0.083333	1.000000	
0.214401			
total_phenols	-0.321113	0.214401	
1.000000			
flavanoids	-0.351370	0.195784	
0.864564			
nonflavanoid_phenols	0.361922	-0.256294	-
0.449935			
proanthocyanins	-0.197327	0.236441	
0.612413			
color_intensity	0.018732	0.199950	-

0.055136		
hue	-0.273955	0.055398
0.433681		
od280/od315_of_diluted_wines	-0.276769	0.066004
0.699949		
proline	-0.440597	0.393351
0.498115		

	flavanoids	nonflavanoid_phenols \
alcohol	0.236815	-0.155929
malic_acid	-0.411007	0.292977
ash	0.115077	0.186230
alcalinity_of_ash	-0.351370	0.361922
magnesium	0.195784	-0.256294
total_phenols	0.864564	-0.449935
flavanoids	1.000000	-0.537900
nonflavanoid_phenols	-0.537900	1.000000
proanthocyanins	0.652692	-0.365845
color_intensity	-0.172379	0.139057
hue	0.543479	-0.262640
od280/od315_of_diluted_wines	0.787194	-0.503270
proline	0.494193	-0.311385


	proanthocyanins	color_intensity
hue \		
alcohol	0.136698	0.546364 -
0.071747		
malic_acid	-0.220746	0.248985 -
0.561296		
ash	0.009652	0.258887 -
0.074667		
alcalinity_of_ash	-0.197327	0.018732 -
0.273955		
magnesium	0.236441	0.199950
0.055398		
total_phenols	0.612413	-0.055136
0.433681		
flavanoids	0.652692	-0.172379
0.543479		
nonflavanoid_phenols	-0.365845	0.139057 -
0.262640		
proanthocyanins	1.000000	-0.025250
0.295544		
color_intensity	-0.025250	1.000000 -
0.521813		
hue	0.295544	-0.521813
1.000000		
od280/od315_of_diluted_wines	0.519067	-0.428815
0.565468		
proline	0.330417	0.316100

0.236183

	od280/od315_of_diluted_wines	proline
alcohol	0.072343	0.643720
malic_acid	-0.368710	-0.192011
ash	0.003911	0.223626
alcalinity_of_ash	-0.276769	-0.440597
magnesium	0.066004	0.393351
total_phenols	0.699949	0.498115
flavanoids	0.787194	0.494193
nonflavanoid_phenols	-0.503270	-0.311385
proanthocyanins	0.519067	0.330417
color_intensity	-0.428815	0.316100
hue	0.565468	0.236183
od280/od315_of_diluted_wines	1.000000	0.312761
proline	0.312761	1.000000

```
import seaborn as sns
import matplotlib.pyplot as plt

# Plot the correlation matrix heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()
```



```

threshold = 0.5
column_names = set()
for row in range(len(corr_matrix)):
    for col in range(row):
        if abs(corr_matrix.iloc[row][col]) > threshold:
            print(corr_matrix.iloc[row][col])
            print(corr_matrix.columns[row])
            column_names.add(corr_matrix.columns[row])

```

```

0.8645635000951147
flavanoids
-0.5378996119051984
nonflavanoid_phenols
0.6124130837800363
proanthocyanins
0.6526917686075153
proanthocyanins
0.546364195083705
color_intensity
-0.5612956886649447

```

```

hue
0.5434785664899897
hue
-0.5218131932287572
hue
0.6999493647911861
od280/od315_of_diluted_wines
0.787193901866951
od280/od315_of_diluted_wines
-0.5032695960789114
od280/od315_of_diluted_wines
0.519067095682523
od280/od315_of_diluted_wines
0.5654682931826589
od280/od315_of_diluted_wines
0.6437200371782134
proline

column_names
{'color_intensity',
 'flavanoids',
 'hue',
 'nonflavanoid_phenols',
 'od280/od315_of_diluted_wines',
 'proanthocyanins',
 'proline'}

```

```
wine_df.head()
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium
total_phenols \					
0	14.23	1.71	2.43	15.6	127.0
2.80					
1	13.20	1.78	2.14	11.2	100.0
2.65					
2	13.16	2.36	2.67	18.6	101.0
2.80					
3	14.37	1.95	2.50	16.8	113.0
3.85					
4	13.24	2.59	2.87	21.0	118.0
2.80					

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity
hue \				
0	3.06	0.28	2.29	5.64
1.04				
1	2.76	0.26	1.28	4.38
1.05				
2	3.24	0.30	2.81	5.68
1.03				

3	3.49	0.24	2.18	7.80
0.86				
4	2.69	0.39	1.82	4.32
1.04				

	od280/od315_of_diluted_wines	proline
0	3.92	1065.0
1	3.40	1050.0
2	3.17	1185.0
3	3.45	1480.0
4	2.93	735.0

```
X2_filt = wine_df.drop(column_names, axis = 1)
```

```
X2_filt.head()
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium
total_phenols					
0	14.23	1.71	2.43	15.6	127.0
2.80					
1	13.20	1.78	2.14	11.2	100.0
2.65					
2	13.16	2.36	2.67	18.6	101.0
2.80					
3	14.37	1.95	2.50	16.8	113.0
3.85					
4	13.24	2.59	2.87	21.0	118.0
2.80					

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.svm import SVC
```

```
y2= wine['target']
```

```
X_train,X_test,y_train,y_test = train_test_split(X2_filt,y2, test_size
=0.24,random_state = 91)
```

```
print(X_train.shape,X_test.shape)
```

```
print(y_train.shape,y_test.shape)
```

```
(135, 6) (43, 6)
```

```
(135,) (43,)
```

```
classifier = SVC(kernel = 'poly')
```

```
classifier.fit(X_train,y_train)
```

```
SVC(kernel='poly')
```

```
y2_pred = classifier.predict(X_test)
```



```
print(" Initial Accuracy - ",accuracy_score(y_test,y2_pred))
print(" Initial f1 score - ",f1_score(y_test,y2_pred, average =
'weighted') )
```

```
Initial Accuracy - 0.5581395348837209
Initial f1 score - 0.4918297498903028
```

Changing Hyperparameters so that we can get better Accuracy

```
from sklearn.model_selection import GridSearchCV
param_grid = {'C': [0.1, 1, 10],
              'kernel': ['linear', 'rbf', 'poly'],
              'gamma': ['scale', 'auto']}
grid_search = GridSearchCV(classifier, param_grid, cv=5)
grid_search.fit(X_train, y_train)
best_svm = grid_search.best_estimator_
best_predictions = best_svm.predict(X_test)
best_accuracy = accuracy_score(y_test, best_predictions)
final_svm = SVC(*grid_search.best_params_)
final_svm.fit(X_train, y_train)
final_predictions = final_svm.predict(X_test)
final_accuracy = accuracy_score(y_test, final_predictions)

print("This is the Best Accuracy after changing Hyper-parameters")
print("Best Hyperparameters:", grid_search.best_params_)
print("Final Accuracy:", final_accuracy)
print(" Final F1
Score:",f1_score(y_test,final_predictions,average="weighted"))
```

```
This is the Best Accuracy after changing Hyper-parameters
Best Hyperparameters: {'C': 10, 'gamma': 'scale', 'kernel': 'linear'}
Final Accuracy: 0.8837209302325582
Final F1 Score: 0.8785295871479182
```

Question NO.3 Using sklearn.datasets.load_diabetes apply Mutual info Classification and check which are the best columns according to the target column. Then Apply decision tree on that data and try to get best accuracy by changing the hyperparameters

```
from sklearn.feature_selection import SelectKBest, mutual_info_classif
from sklearn.datasets import load_diabetes
from sklearn.model_selection import GridSearchCV
```

```
data = load_diabetes()
X3 = data.data
y3 = data.target
```

```
df = pd.DataFrame(X3, columns=data.feature_names)
df['target'] = y3
```

```
df.head()
```

```

      age      sex      bmi      bp      s1      s2
s3  \
```

```

0  0.038076  0.050680  0.061696  0.021872 -0.044223 -0.034821 -
0.043401
1 -0.001882 -0.044642 -0.051474 -0.026328 -0.008449 -0.019163
0.074412
2  0.085299  0.050680  0.044451 -0.005671 -0.045599 -0.034194 -
0.032356
3 -0.089063 -0.044642 -0.011595 -0.036656  0.012191  0.024991 -
0.036038
4  0.005383 -0.044642 -0.036385  0.021872  0.003935  0.015596
0.008142

```

```

          s4          s5          s6  target
0 -0.002592  0.019908 -0.017646   151.0
1 -0.039493 -0.068330 -0.092204    75.0
2 -0.002592  0.002864 -0.025930   141.0
3  0.034309  0.022692 -0.009362   206.0
4 -0.002592 -0.031991 -0.046641   135.0

```

```

selector = SelectKBest(score_func=mutual_info_classif, k=4)
X_new = selector.fit_transform(X3, y3)

```

```

selected_features =
df.columns[selector.get_support(indices=True)].tolist()

```

```

print("Selected Features:", selected_features)

```

```

Selected Features: ['sex', 's3', 's4', 's5']

```

```

X_filt = df[['sex', 's3', 's4', 's5']]
X_filt.head()

```

```

          sex          s3          s4          s5
0  0.050680 -0.043401 -0.002592  0.019908
1 -0.044642  0.074412 -0.039493 -0.068330
2  0.050680 -0.032356 -0.002592  0.002864
3 -0.044642 -0.036038  0.034309  0.022692
4 -0.044642  0.008142 -0.002592 -0.031991

```

```

df['target'] = y3

```

```

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import r2_score
X_train,X_test,y_train,y_test = train_test_split(X_filt,y3, test_size
=0.29,random_state = 17)
print(X_train.shape,X_test.shape)
print(y_train.shape,y_test.shape)
tree = DecisionTreeClassifier(criterion = "entropy")
tree.fit(X_train,y_train)

```

```

y_pred = tree.predict(X_test)
print("Accuracy of Decision Tree Model
",accuracy_score(y_test,y_pred))

print("r2 ",r2_score(y_test,y_pred))
print(" Initial Accuracy - ",accuracy_score(y_test,y_pred))
print(" Initial f1 score - ",f1_score(y_test,y_pred, average =
'weighted') )

```

```

(313, 4) (129, 4)
(313,) (129,)
Accuracy of Decision Tree Model  0.015503875968992248
r2  -0.3063360751715225
Initial Accuracy -  0.015503875968992248
Initial f1 score -  0.012919896640826873

```

Define the hyperparameters to optimize

```

hyperparameters = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [5, 8, 10],
    'min_samples_split': [2, 4, 8],
    'min_samples_leaf': [1, 2, 4]
}

```

```

grid_search = GridSearchCV(tree, hyperparameters, cv=2)
grid_search.fit(X_train, y_train)

```

```

best_model = grid_search.best_estimator_

```

```

y_pred = best_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

```

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\
_split.py:676: UserWarning: The least populated class in y has only 1
members, which is less than n_splits=2.

```

```

    warnings.warn(

```

```

print("Best columns:", selected_features)
print("Best hyperparameters:", grid_search.best_params_)
print("Best accuracy:", accuracy)

```

```

Best columns: ['sex', 's3', 's4', 's5']
Best hyperparameters: {'criterion': 'gini', 'max_depth': 5,
'min_samples_leaf': 4, 'min_samples_split': 8}
Best accuracy: 0.007751937984496124

```

Question NO. 4 Using `sklearn.datasets.load_boston` apply Mutual info Regression and check which are the best columns according to the target column. Then Apply MultiLinear Regression on that data and try to get best accuracy by changing the hyperparameters

```
from sklearn.feature_selection import SelectKBest,  
mutual_info_regression  
from sklearn.datasets import load_boston
```

```
bos = load_boston()
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\  
deprecation.py:87: FutureWarning: Function load_boston is deprecated;  
`load_boston` is deprecated in 1.0 and will be removed in 1.2.
```

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd  
import numpy as np  
  
data_url = "http://lib.stat.cmu.edu/datasets/boston"  
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22,  
header=None)  
data = np.hstack([raw_df.values[::2, :],  
raw_df.values[1::2, :2]])  
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch_california_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing  
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
```

```

housing = fetch_openml(name="house_prices", as_frame=True)

for the Ames housing dataset.

warnings.warn(msg, category=FutureWarning)

bos.feature_names
array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',
      'RAD',
      'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')

X4 = bos.data
y4= bos.target
bos_data = pd.DataFrame(X4, columns=bos.feature_names)
bos_data.head()

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

```

bos_data['target'] = y4

selector = SelectKBest(score_func=mutual_info_regression, k=4)
X_new = selector.fit_transform(X4, y4)
selected_features =
bos_data.columns[selector.get_support(indices=True)].tolist()
print("Selected Features:", selected_features)

Selected Features: ['INDUS', 'NOX', 'RM', 'LSTAT']

X_filt1 = bos_data[['INDUS', 'NOX', 'RM', 'LSTAT']]
X_filt1.head()

```

	INDUS	NOX	RM	LSTAT
0	2.31	0.538	6.575	4.98
1	7.07	0.469	6.421	9.14
2	7.07	0.469	7.185	4.03
3	2.18	0.458	6.998	2.94
4	2.18	0.458	7.147	5.33

```
X_train,X_test,y_train,y_test = train_test_split(X_filt1,y4, test_size
=0.28,random_state = 63)
```

```
print(X_train.shape,X_test.shape)
```

```
print(y_train.shape,y_test.shape)
```

```
from sklearn import linear_model
```

```
model= linear_model.LinearRegression()
```

```
model.fit(X_train,y_train)
```

```
y_pred = model.predict(X_test)
```

```
print("Accuracy for this model",r2_score(y_test,y_pred))
```

```
(364, 4) (142, 4)
```

```
(364,) (142,)
```

```
Accuracy for this model 0.628740044008688
```

```
#change the hyperparameters to get better accuracy
```

```
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.pipeline import make_pipeline
```

```
from sklearn.preprocessing import StandardScaler
```

```
param_grid = {
```

```
    'linearregression__normalize': [True, False],
```

```
    'linearregression__fit_intercept': [True, False]
```

```
}
```

```
pipeline = make_pipeline(StandardScaler(), LinearRegression())
```

```
grid_search = GridSearchCV(pipeline, param_grid, cv=5,
```

```
scoring='neg_mean_squared_error')
```

```
grid_search.fit(X_train, y_train)
```

```
best_params = grid_search.best_params_
```

```
best_model = grid_search.best_estimator_
```

```
y_pred_best = best_model.predict(X_test)
```

```
r2_best = r2_score(y_test, y_pred_best)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
```

```
If you wish to scale the data, use Pipeline with a StandardScaler in a
preprocessing stage. To reproduce the previous behavior:
```

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
```

```
LinearRegression())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
If you wish to scale the data, use Pipeline with a StandardScaler in a
preprocessing stage. To reproduce the previous behavior:
```

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
LinearRegression())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
If you wish to scale the data, use Pipeline with a StandardScaler in a
preprocessing stage. To reproduce the previous behavior:
```

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
LinearRegression())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
```

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
LinearRegression())
```

If you wish to pass a sample_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
```

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
LinearRegression())
```

If you wish to pass a sample_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
```


and will be removed in 1.2. Please leave the normalize parameter to its default value to silence this warning. The default behavior of this estimator is to not do any normalization. If normalization is needed please use `sklearn.preprocessing.StandardScaler` instead.

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
```

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
LinearRegression())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
```

If you wish to scale the data, use Pipeline with a StandardScaler in a

preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),  
LinearRegression())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in  
model.steps}  
model.fit(X, y, **kwargs)
```

```
warnings.warn(  
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\  
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0  
and will be removed in 1.2.
```

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),  
LinearRegression())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in  
model.steps}  
model.fit(X, y, **kwargs)
```

```
warnings.warn(  
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\  
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0  
and will be removed in 1.2.
```

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),  
LinearRegression())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
If you wish to scale the data, use Pipeline with a StandardScaler in a
preprocessing stage. To reproduce the previous behavior:
```

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
LinearRegression())
```

If you wish to pass a sample_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
```

```
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:148: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2. Please leave the normalize parameter to
its default value to silence this warning. The default behavior of
this estimator is to not do any normalization. If normalization is
needed please use sklearn.preprocessing.StandardScaler instead.
```

```
warnings.warn(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\
_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0
and will be removed in 1.2.
```

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False),
LinearRegression())
```

If you wish to pass a sample_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in
model.steps}
model.fit(X, y, **kwargs)
```

```
warnings.warn(
```

```
print("Best R^2 Score:", r2_best)
print("Best Model Parameters:", best_params)
```

```
Best R^2 Score: 0.628740044008688
```

```
Best Model Parameters: {'linearregression__fit_intercept': True,
'linearregression__normalize': True}
```